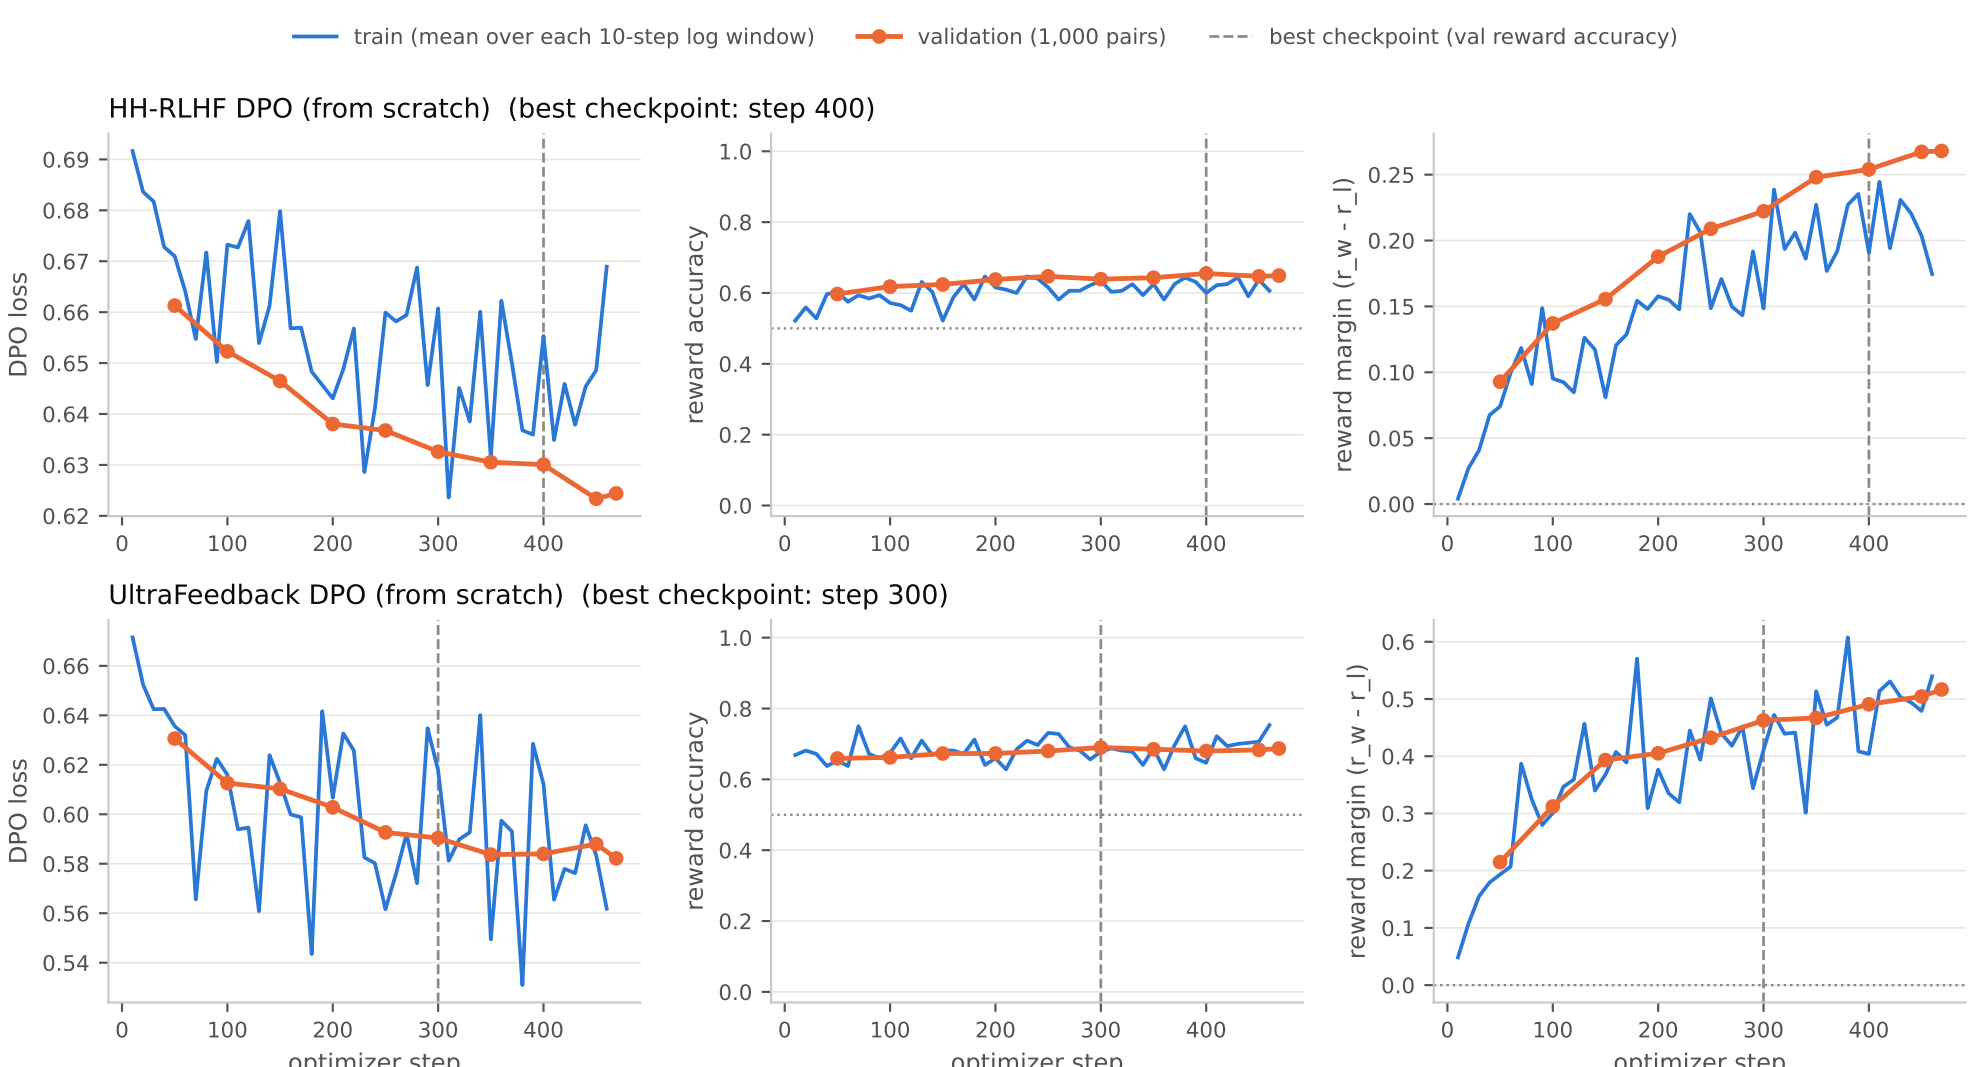
train (mean over each 10-step log window)
validation (1,000 pairs)
best checkpoint (val reward accuracy)
HH-RLHF DPO (from scratch) (best checkpoint: step 400)
DPO loss
reward accuracy
reward margin (r_w - r_l)
UltraFeedback DPO (from scratch) (best checkpoint: step 300)
optimizer step