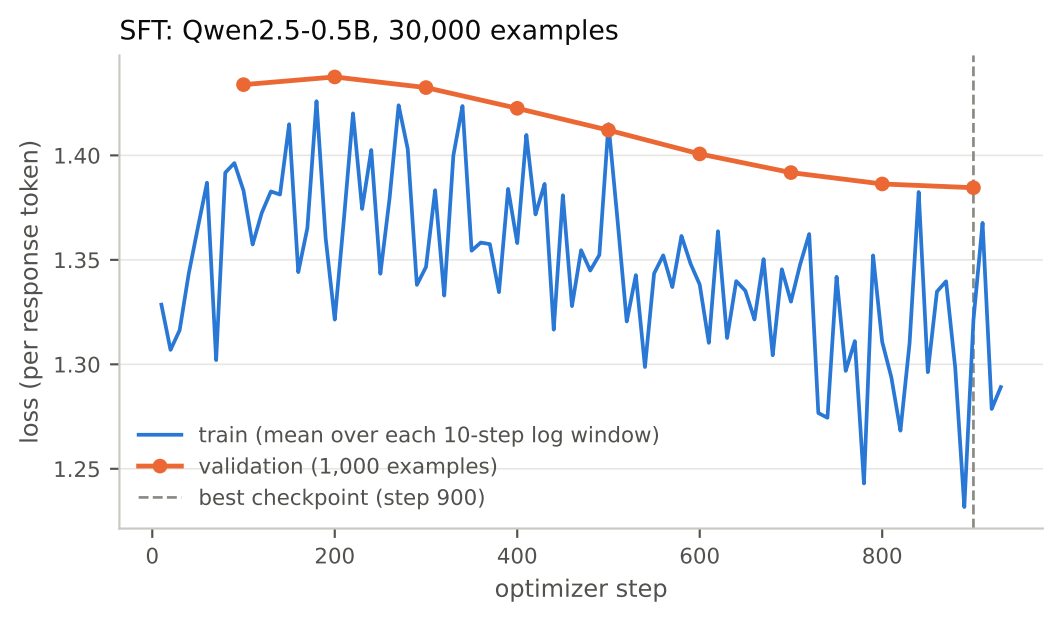
SFT: Qwen2.5-0.5B, 30,000 examples
loss (per response token)
1.40
1.35
1.30
1.25
0
200
400
600
800
optimizer step
train (mean over each 10-step log window)
validation (1,000 examples)
best checkpoint (step 900)