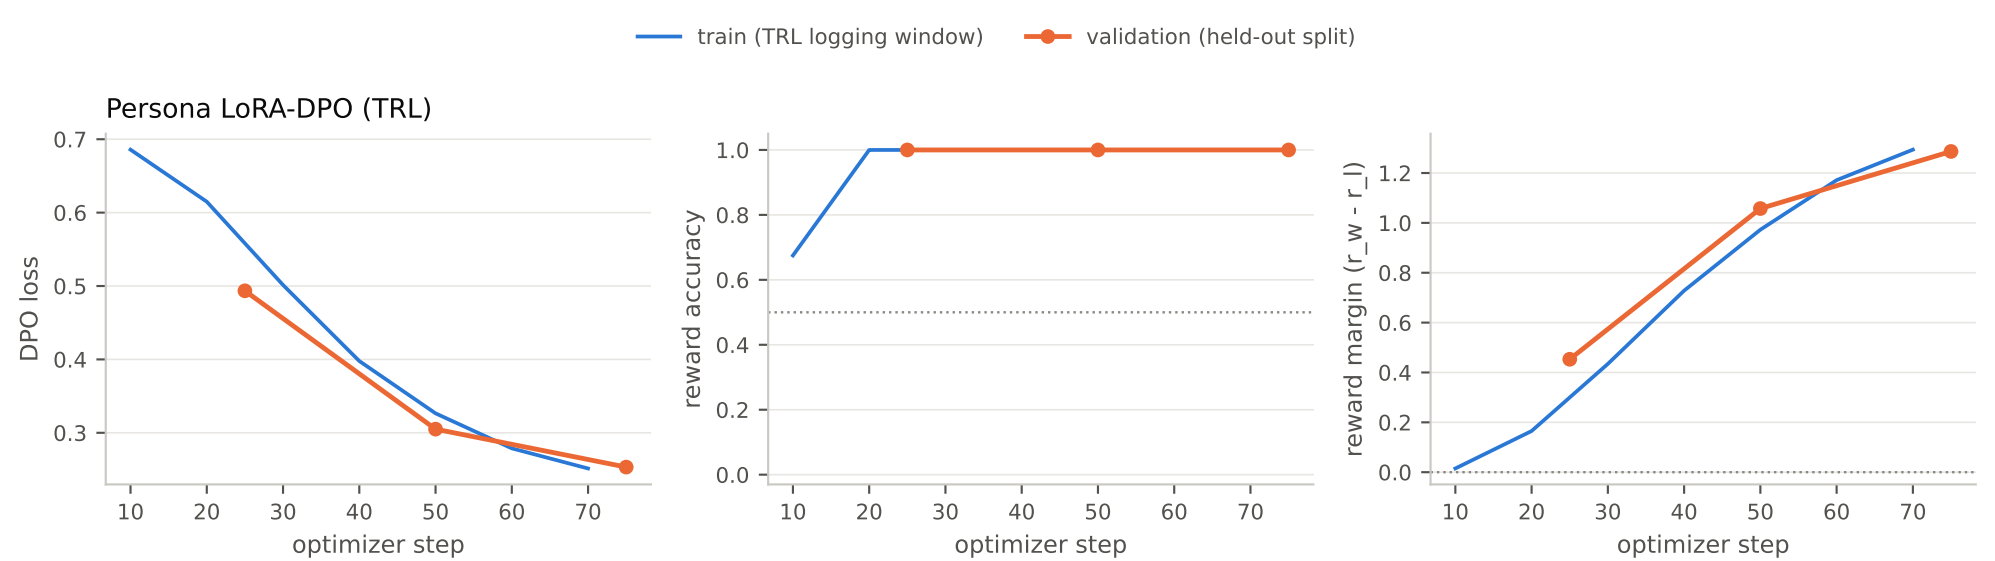
train (TRL logging window)
validation (held-out split)
Persona LoRA-DPO (TRL)
DPO loss
0.7
0.6
0.5
0.4
0.3
10
20
30
40
50
60
70
optimizer step
reward accuracy
1.0
0.8
0.6
0.4
0.2
0.0
10
20
30
40
50
60
70
optimizer step
reward margin (r_w - r_l)
1.2
1.0
0.8
0.6
0.4
0.2
0.0
10
20
30
40
50
60
70
optimizer step